

AI Job Market Visualization

- Explore salaries, skills, remote work, and trends in the global AI job market.
- **Dataset Time Period:** January 2024 - May 2025

Step 0: Import libraries and load data

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
import numpy as np
import warnings
warnings.filterwarnings('ignore')
# For better plots
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (10,6)
# Load dataset
df = pd.read_csv('aijob.csv')
# Quick peek
print(df.head())
print(f"Dataset shape: {df.shape}")
```

	job_id	job_title	salary_usd	salary_currency	\
0	AI00001	AI Research Scientist	90376	USD	
1	AI00002	AI Software Engineer	61895	USD	
2	AI00003	AI Specialist	152626	USD	
3	AI00004	NLP Engineer	80215	USD	
4	AI00005	AI Consultant	54624	EUR	

	experience_level	employment_type	company_location	company_size	\
0	SE	CT	China	M	
1	EN	CT	Canada	M	
2	MI	FL	Switzerland	L	
3	SE	FL	India	M	
4	EN	PT	France	S	

	employee_residence	remote_ratio	\
0	China	50	
1	Ireland	100	
2	South Korea	0	
3	India	50	
4	Singapore	100	

	required_skills	education_required
0	Tableau, PyTorch, Kubernetes, Linux, NLP	Bachelor
1	Deep Learning, AWS, Mathematics, Python, Docker	Master

2	Kubernetes, Deep Learning, Java, Hadoop, NLP	Associate
3	Scala, SQL, Linux, Python	PhD
4	MLOps, Java, Tableau, Python	Master

	years_experience	industry	posting_date	application_deadline	\
0	9	Automotive	2024-10-18	2024-11-07	
1	1	Media	2024-11-20	2025-01-11	
2	2	Education	2025-03-18	2025-04-07	
3	7	Consulting	2024-12-23	2025-02-24	
4	0	Media	2025-04-15	2025-06-23	

	job_description_length	benefits_score	company_name
0	1076	5.9	Smart Analytics
1	1268	5.2	TechCorp Inc
2	1974	9.4	Autonomous Tech
3	1345	8.6	Future Systems
4	1989	6.6	Advanced Robotics

Dataset shape: (15000, 19)

Step 1: Basic info & cleaning

```
print(df.info())
```

```
print(df.isnull().sum())
```

Drop columns if you want, or fill missing values for key cols
For example, fill missing salary with median or drop rows with missing salary

```
df = df.dropna(subset=['salary_usd', 'experience_level', 'industry'])
```

Convert posting_date to datetime

```
df['posting_date'] = pd.to_datetime(df['posting_date'],
errors='coerce')
```

Clean remote_ratio to numeric if needed

```
df['remote_ratio'] = pd.to_numeric(df['remote_ratio'],
errors='coerce')
```

Quick check

```
print(df.describe())
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 15000 entries, 0 to 14999
```

```
Data columns (total 19 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	job_id	15000 non-null	object
1	job_title	15000 non-null	object
2	salary_usd	15000 non-null	int64
3	salary_currency	15000 non-null	object
4	experience_level	15000 non-null	object
5	employment_type	15000 non-null	object
6	company_location	15000 non-null	object

7	company_size	15000	non-null	object
8	employee_residence	15000	non-null	object
9	remote_ratio	15000	non-null	int64
10	required_skills	15000	non-null	object
11	education_required	15000	non-null	object
12	years_experience	15000	non-null	int64
13	industry	15000	non-null	object
14	posting_date	15000	non-null	object
15	application_deadline	15000	non-null	object
16	job_description_length	15000	non-null	int64
17	benefits_score	15000	non-null	float64
18	company_name	15000	non-null	object

dtypes: float64(1), int64(4), object(14)

memory usage: 2.2+ MB

None

job_id	0
job_title	0
salary_usd	0
salary_currency	0
experience_level	0
employment_type	0
company_location	0
company_size	0
employee_residence	0
remote_ratio	0
required_skills	0
education_required	0
years_experience	0
industry	0
posting_date	0
application_deadline	0
job_description_length	0
benefits_score	0
company_name	0

dtype: int64

	salary_usd	remote_ratio	years_experience \
count	15000.000000	15000.000000	15000.000000
mean	115348.965133	49.483333	6.253200
min	32519.000000	0.000000	0.000000
25%	70179.750000	0.000000	2.000000
50%	99705.000000	50.000000	5.000000
75%	146408.500000	100.000000	10.000000
max	399095.000000	100.000000	19.000000
std	60260.940438	40.812712	5.545768

	posting_date	job_description_length
benefits_score		
count	15000	15000.000000
15000.000000		

mean	2024-08-29 08:48:51.840000	1503.314733
7.504273		
min	2024-01-01 00:00:00	500.000000
5.000000		
25%	2024-04-29 00:00:00	1003.750000
6.200000		
50%	2024-08-28 00:00:00	1512.000000
7.500000		
75%	2024-12-29 00:00:00	2000.000000
8.800000		
max	2025-04-30 00:00:00	2499.000000
10.000000		
std	NaN	576.127083
1.450870		

Salary Distribution by Experience Level

```
# Step 2: Salary vs Experience Level
plt.figure(figsize=(8,5))
sns.boxplot(data=df, x='experience_level', y='salary_usd',
order=['EN', 'MI', 'SE', 'EX'])
plt.title("Salary Distribution by Experience Level")
plt.ylabel("Salary (USD)")
plt.xlabel("Experience Level")
plt.show()
```



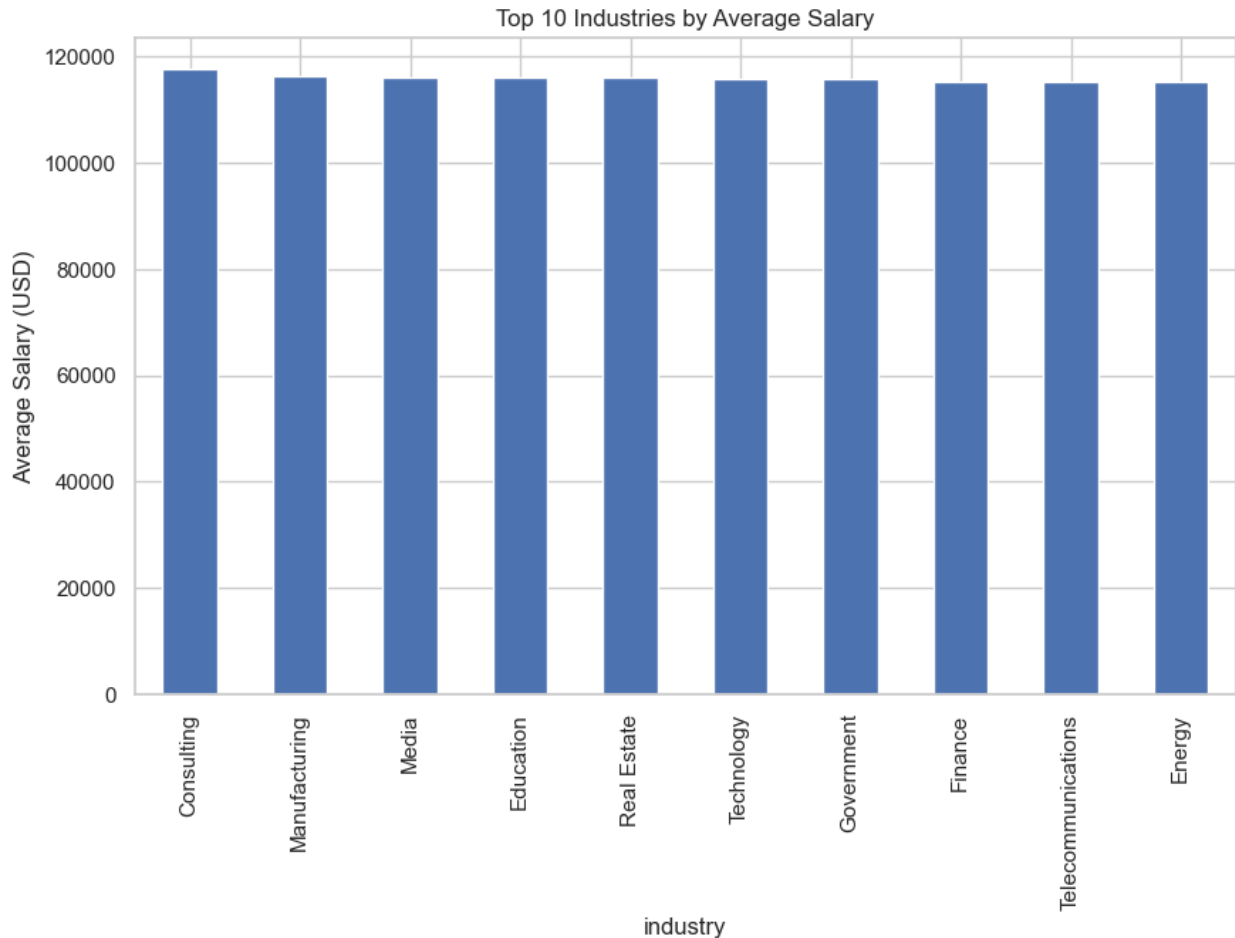
Average Salary vs Years of Experience

```
# Step 3: Scatter plot salary vs years_experience
plt.figure(figsize=(8,5))
sns.lineplot(data=df, x='years_experience', y='salary_usd',
             estimator='mean')
plt.title("Average Salary vs Years of Experience")
plt.xlabel("Years of Experience")
plt.ylabel("Salary (USD)")
plt.show()
```



Top 10 Industries by Average Salary


```
# Step 4: Top 10 industries by average salary
industry_salary = df.groupby('industry')
['salary_usd'].mean().sort_values(ascending=False).head(10)
industry_salary.plot(kind='bar')
plt.title("Top 10 Industries by Average Salary")
plt.ylabel("Average Salary (USD)")
plt.show()
```

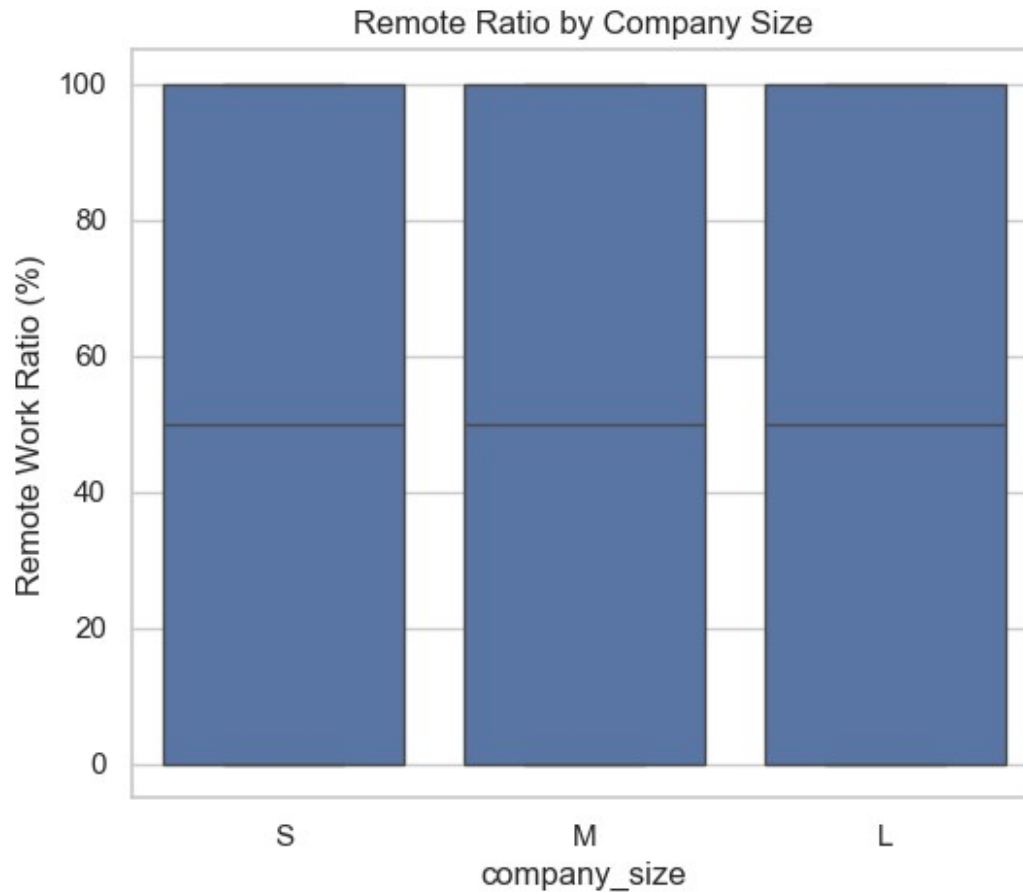


Remote Work Ratio by Industry and Remote Work by Company Size

```
# Step 5: Remote ratio distribution by industry
plt.figure(figsize=(12,6))
sns.boxplot(data=df, x='industry', y='remote_ratio')
plt.xticks(rotation=90)
plt.title("Remote Ratio by Industry")
plt.ylabel("Remote Work Ratio (%)")
plt.show()

# Remote ratio by company size
plt.figure(figsize=(6,5))
sns.boxplot(data=df, x='company_size', y='remote_ratio', order=['S', 'M', 'L'])
plt.title("Remote Ratio by Company Size")
plt.ylabel("Remote Work Ratio (%)")
plt.show()
```






On-site vs Remote Jobs

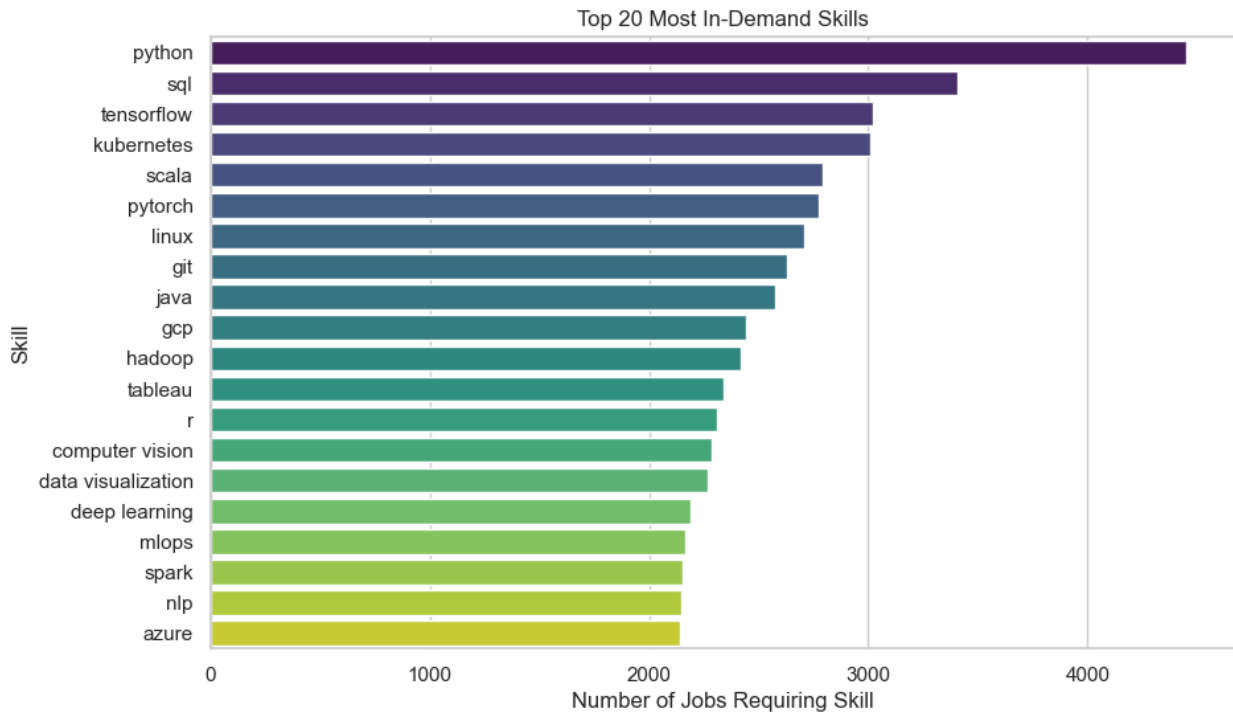
```
# Step 6: Count jobs where employee_residence == company_location (on-site) vs remote
df['is_remote'] = np.where(df['employee_residence'] ==
df['company_location'], 'On-site', 'Remote')
remote_counts = df['is_remote'].value_counts()
remote_counts.plot(kind='pie', autopct='%1.1f%%', startangle=90,
colors=['#66b3ff','#ff9999'])
plt.title("On-site vs Remote Jobs")
plt.ylabel('')
plt.show()
```

On-site vs Remote Jobs



Top 20 Most In-Demand Skills

```
# Step 7: Extract and count required skills
# Assuming skills are comma separated strings
all_skills =
df['required_skills'].dropna().str.split(',').explode().str.strip().str.lower()
skill_counts = Counter(all_skills)
# Top 20 skills
top_skills = pd.DataFrame(skill_counts.most_common(20),
columns=['Skill', 'Count'])
sns.barplot(data=top_skills, x='Count', y='Skill', palette='viridis')
plt.title("Top 20 Most In-Demand Skills")
plt.xlabel("Number of Jobs Requiring Skill")
plt.show()
```

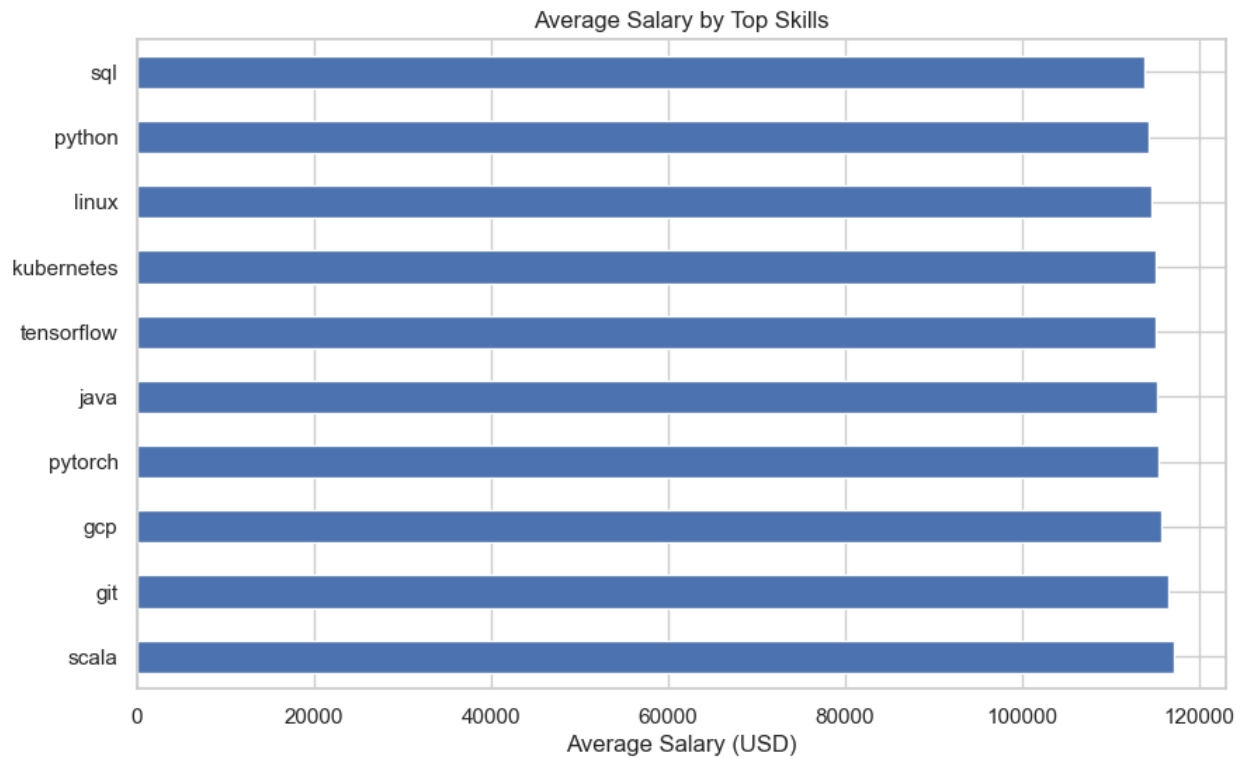


Average Salary for Top 10 Skills

Step 8: Average salary for top 10 skills


Filter dataset to rows that contain each top skill and calculate average salary

```
avg_salary_skills = {}
for skill in top_skills['Skill'][:10]:
    mask = df['required_skills'].str.lower().str.contains(skill,
na=False)
    avg_salary_skills[skill] = df.loc[mask, 'salary_usd'].mean()
avg_salary_df = pd.DataFrame.from_dict(avg_salary_skills,
orient='index', columns=['Average_Salary'])
avg_salary_df = avg_salary_df.sort_values(by='Average_Salary',
ascending=False)
avg_salary_df.plot(kind='barh', legend=False)
plt.title("Average Salary by Top Skills")
plt.xlabel("Average Salary (USD)")
plt.show()
```



Jobs Posted Over Time

```
# Step 9: Number of jobs posted monthly
df['posting_month'] = df['posting_date'].dt.to_period('M')
jobs_per_month = df.groupby('posting_month').size()
jobs_per_month.plot()
plt.title("Number of Jobs Posted Over Time")
plt.ylabel("Number of Job Postings")
plt.xlabel("Month")
plt.show()
```



Salary & Benefits by Company Size

```
salary_by_size = df.groupby('company_size')
['salary_usd'].mean().reindex(['S', 'M', 'L'])
benefits_by_size = df.groupby('company_size')
['benefits_score'].mean().reindex(['S', 'M', 'L'])

# Plot
fig, ax = plt.subplots(1, 2, figsize=(14, 5))
sns.set_theme(style="whitegrid", font_scale=1.1)

# Bar plot: Salary
sns.barplot(x=salary_by_size.index, y=salary_by_size.values,
palette='Blues_d', ax=ax[0])
ax[0].set_title("Average Salary by Company Size", fontweight='bold')
ax[0].set_xlabel("Company Size")
ax[0].set_ylabel("Average Salary (USD)")
ax[0].bar_label(ax[0].containers[0], fmt='%.0f', padding=3)

# Bar plot: Benefits Score
sns.barplot(x=benefits_by_size.index, y=benefits_by_size.values,
palette='Greens_d', ax=ax[1])
ax[1].set_title("Average Benefits Score by Company Size",
fontweight='bold')
```

```

ax[1].set_xlabel("Company Size")
ax[1].set_ylabel("Average Benefits Score")
ax[1].bar_label(ax[1].containers[0], fmt='%.2f', padding=3)

plt.suptitle("Comparison of Salary & Benefits Across Company Sizes",
fontsize=15, fontweight='bold', y=1.03)
plt.tight_layout()
plt.show()

```

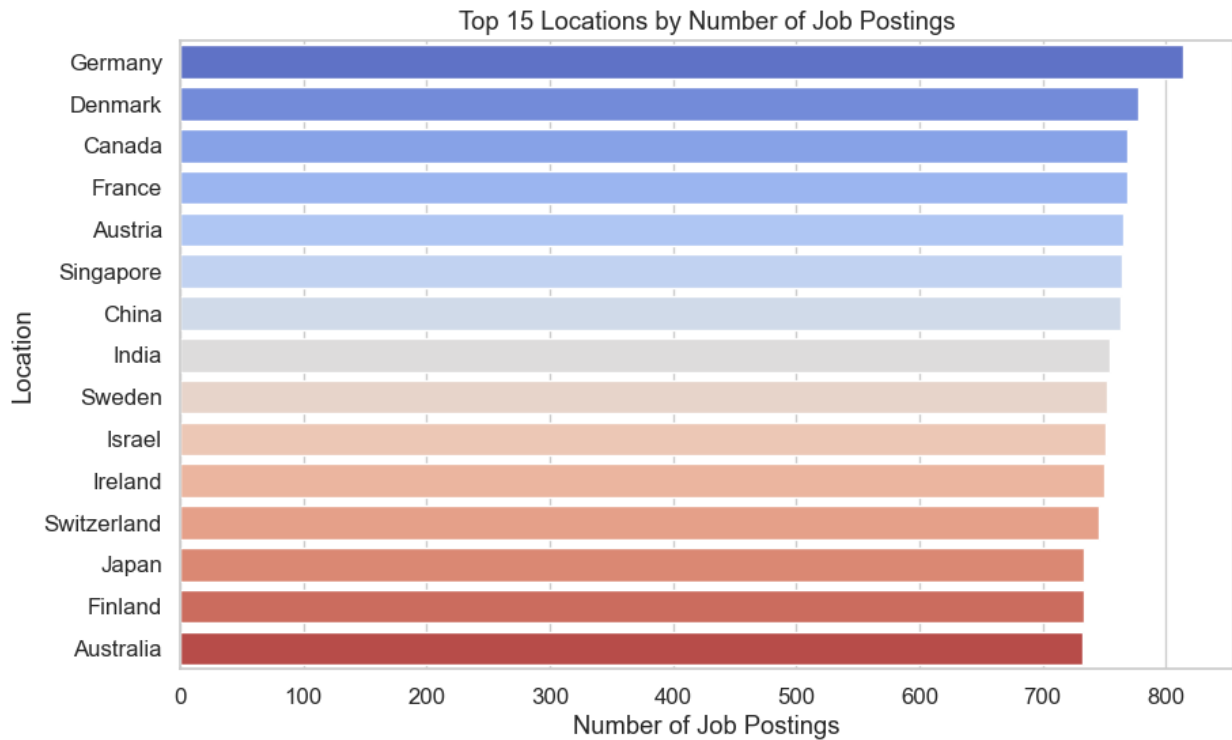


Top 15 Job Locations

```

top_locations = df['company_location'].value_counts().head(15)
plt.figure(figsize=(10,6))
sns.barplot(x=top_locations.values, y=top_locations.index,
palette='coolwarm')
plt.title('Top 15 Locations by Number of Job Postings')
plt.xlabel('Number of Job Postings')
plt.ylabel('Location')
plt.show()


```



Top 15 Locations by Salary

```
avg_salary_loc = df.groupby('company_location')
['salary_usd'].mean().sort_values(ascending=False).head(15)

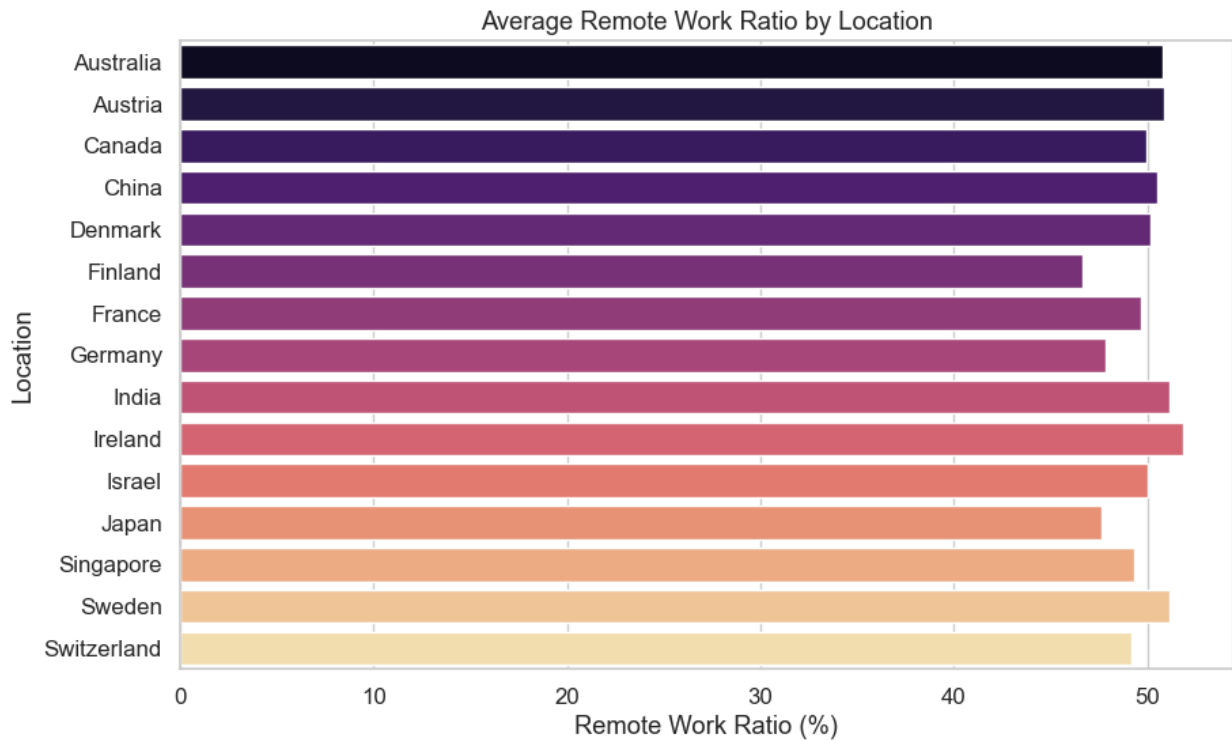
plt.figure(figsize=(10,6))
sns.barplot(x=avg_salary_loc.values, y=avg_salary_loc.index,
palette='viridis')
plt.title('Top 15 Locations by Average Salary (USD)')
plt.xlabel('Average Salary')
plt.ylabel('Location')
plt.show()
```



Percentage of Remote Work Done by Employees in Various Countries

```
# Calculate average remote_ratio per location (only locations with enough data)
remote_loc = df.groupby('company_location')['remote_ratio'].mean()
remote_loc = remote_loc[remote_loc.index.isin(top_locations.index)]

plt.figure(figsize=(10,6))
sns.barplot(x=remote_loc.values, y=remote_loc.index, palette='magma')
plt.title('Average Remote Work Ratio by Location')
plt.xlabel('Remote Work Ratio (%)')
plt.ylabel('Location')
plt.show()
```


Job Posting Trend Over Time (Top 5 Locations)

```
# Filter for top 5 locations for clarity
top5_locations = top_locations.index[:5]

df['posting_month'] = df['posting_date'].dt.to_period('M')

jobs_time_loc =
df[df['company_location'].isin(top5_locations)].groupby(['posting_month', 'company_location']).size().unstack(fill_value=0)

jobs_time_loc.plot(figsize=(12,6), marker='o')
plt.title('Job Posting Trend Over Time (Top 5 Locations)')
plt.xlabel('Month')
plt.ylabel('Number of Job Postings')
plt.legend(title='Location')
plt.show()
```



Salary Heatmap (Top 10 Locations × Top 8 Industries)

```
# Top 10 locations and top 8 industries
top10_locations = df['company_location'].value_counts().head(10).index
top8_industries = df['industry'].value_counts().head(8).index

pivot_salary = df[(df['company_location'].isin(top10_locations)) &
                  (df['industry'].isin(top8_industries))]\
    .pivot_table(values='salary_usd', index='company_location',
                  columns='industry', aggfunc='mean')

plt.figure(figsize=(12,7))
sns.heatmap(pivot_salary, annot=True, fmt=".0f", cmap='YlGnBu')
plt.title('Average Salary by Location and Industry')
plt.xlabel('Industry')
plt.ylabel('Location')
plt.show()
```

